

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib import rcParams
rcParams.update({'figure.autolayout': True})
import seaborn as sns
import joblib
from sklearn.feature_selection import SelectFromModel
from sklearn.linear_model import LassoCV
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import FunctionTransformer, OneHotEncoder, LabelEncoder
from sklearn.utils.multiclass import type_of_target
from sklearn.utils import check_X_y
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score, RepeatedStratifiedKFold, KFold
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, roc_auc_score
%run './python/Chicago_Car_Crashes.py'
```

DATA COLLECTION

```
In [103]: crashes = pd.read_csv('./data/Traffic_Crashes_-_Crashes.csv.gz', index_col=0)
print(f"{'Dataset contains shape'}\n{'Rows':10}{crashes.shape[0]}\n{'Columns':10}{crashes.shape[1]}")
```

Dataset contains shape
Rows 490128
Columns 49

```
In [4]: crashes.CRASH_DATE = pd.to_datetime(crashes.CRASH_DATE)
df = crashes[(crashes.CRASH_DATE.dt.year == 2020)].drop('Unnamed: 0', axis=1).copy()
df.to_csv('./data/Chicago_Traffic_Crashes_2020.csv.gz', index=False, compression='gzip')
```

```
In [2]: df = pd.read_csv('./data/Chicago_Traffic_Crashes_2020.csv.gz', low_memory=False)
```

DESCRIPTIVE STATISTICS

```
In [3]: df.head()
```

	CRASH_RECORD_ID	RD_NO	CRASH_DATE_EST_I	CRASH_DATE	POSTED_SPEED_LIMIT	TRAFFIC_CONTROL
0	ee9283eff3a55ac50ee58f3d9528ce1d689b1c4180b4c4...	JD292400	NaN	2020-07-10 10:25:00	30	TRAFFIC
1	f8960f698e870ebdc60b521b2a141a5395556bc3704191...	JD293602	NaN	2020-07-11 01:00:00	30	NO CC
2	8eaa2678d1a127804ee9b8c35ddf7d63d913c14eda61d6...	JD290451	NaN	2020-07-08 14:00:00	20	NO CC
3	f636d4a51a88015ac89031159b1f1952b8d92e49d11aeb...	JD293237	NaN	2020-07-10 22:20:00	30	NO CC
4	76aabcf7c2219a5c90259c96fe94b33834ddb53e0dbcd3...	JD291826	NaN	2020-07-09 17:06:00	10	NO CC

5 rows x 49 columns

```
In [4]: df.describe()
```

	POSTED_SPEED_LIMIT	LANE_CNT	STREET_NO	BEAT_OF_OCCURRENCE	NUM_UNITS	INJURIES_TOTAL	INJURIES_FATAL	INJURIES_INC/
count	92092.000000	12.000000	92092.000000	92092.000000	92092.00000	91858.000000	91858.000000	9
mean	28.532793	2.250000	3860.405985	1216.832732	2.04831	0.215060	0.001698	
std	5.761715	0.965307	2905.455298	697.276831	0.50948	0.600362	0.044480	
min	0.000000	1.000000	0.000000	111.000000	1.00000	0.000000	0.000000	
25%	30.000000	2.000000	1401.000000	633.000000	2.00000	0.000000	0.000000	
50%	30.000000	2.000000	3407.000000	1121.000000	2.00000	0.000000	0.000000	
75%	30.000000	2.250000	5749.000000	1731.000000	2.00000	0.000000	0.000000	
max	70.000000	4.000000	13799.000000	6100.000000	15.00000	19.000000	3.000000	

DATA SCRUBBING

prima di iniziare ad esplorare i dati dobbiamo individuare le variable da escludere in particolare quelli con i dati mancanti e quelli che hanno poca rilevanza per il nostro modello.

```
In [5]: dfi = Dframe(df)
```

```
In [6]: dfi.df_info()
```

Data Entries	92092
Data Columns	49
DataFrame items	4,512,508
DataFrame Null	948953 Null values
DataFrame contain	21.03% Null Values

#	Column	Null Count &	Percent	Dtype
---	-----	-----	-----	-----
0	CRASH_RECORD_ID	0	0.00	% object
1	RD_NO	0	0.00	% object
2	CRASH_DATE_EST_I	85,157	92.47	% object
3	CRASH_DATE	0	0.00	% object
4	POSTED_SPEED_LIMIT	0	0.00	% int64
5	TRAFFIC_CONTROL_DEVICE	0	0.00	% object
6	DEVICE_CONDITION	0	0.00	% object
7	WEATHER_CONDITION	0	0.00	% object
8	LIGHTING_CONDITION	0	0.00	% object
9	FIRST_CRASH_TYPE	0	0.00	% object
10	TRAFFICWAY_TYPE	0	0.00	% object
11	LANE_CNT	92,080	99.99	% float64
12	ALIGNMENT	0	0.00	% object
13	ROADWAY_SURFACE_COND	0	0.00	% object
14	ROAD_DEFECT	0	0.00	% object
15	REPORT_TYPE	2,623	2.85	% object
16	CRASH_TYPE	0	0.00	% object
17	INTERSECTION_RELATED_I	69,300	75.25	% object
18	NOT_RIGHT_OF_WAY_I	87,489	95.00	% object
19	HIT_AND_RUN_I	60,165	65.33	% object
20	DAMAGE	0	0.00	% object
21	DATE_POLICE_NOTIFIED	0	0.00	% object
22	PRIM_CONTRIBUTORY_CAUSE	0	0.00	% object
23	SEC_CONTRIBUTORY_CAUSE	0	0.00	% object
24	STREET_NO	0	0.00	% int64
25	STREET_DIRECTION	0	0.00	% object
26	STREET_NAME	0	0.00	% object
27	BEAT_OF_OCCURRENCE	0	0.00	% float64
28	PHOTOS_TAKEN_I	90,961	98.77	% object
29	STATEMENTS_TAKEN_I	90,310	98.06	% object
30	DOORING_I	91,834	99.72	% object
31	WORK_ZONE_I	91,662	99.53	% object
32	WORK_ZONE_TYPE	91,760	99.64	% object
33	WORKERS_PRESENT_I	91,975	99.87	% object
34	NUM_UNITS	0	0.00	% int64
35	MOST_SEVERE_INJURY	235	0.26	% object
36	INJURIES_TOTAL	234	0.25	% float64
37	INJURIES_FATAL	234	0.25	% float64
38	INJURIES_INCAPACITATING	234	0.25	% float64
39	INJURIES_NON_INCAPACITATING	234	0.25	% float64
40	INJURIES_REPORTED_NOT_EVIDENT	234	0.25	% float64
41	INJURIES_NO_INDICATION	234	0.25	% float64
42	INJURIES_UNKNOWN	234	0.25	% float64
43	CRASH_HOUR	0	0.00	% int64
44	CRASH_DAY_OF_WEEK	0	0.00	% int64
45	CRASH_MONTH	0	0.00	% int64
46	LATITUDE	588	0.64	% float64
47	LONGITUDE	588	0.64	% float64
48	LOCATION	588	0.64	% object

In [7]:

```
to_drop_cols      = set()
excluded_cols     = set()
category_cols     = set()
numeric_cols      = set()
encoding_cols     = set()
features_cols     = set()
statistic_cols    = set()
```

In [8]:

```
excluded_cols.update([
    'Unnamed: 0', 'Unnamed: 0.1', 'CRASH_RECORD_ID', 'RD_NO', 'CRASH_DATE_EST_I', 'LANE_CNT',
    'REPORT_TYPE', 'DAMAGE', 'PHOTOS_TAKEN_I', 'STATEMENTS_TAKEN_I', 'LOCATION',
    'WORK_ZONE_TYPE', 'BEAT_OF_OCCURRENCE', 'DATE_POLICE_NOTIFIED', 'CRASH_TYPE', 'SEC_CONTRIBUTORY_CAUSE'])
```

In [9]:

```
TransformerDate = DateTransformer()
df.CRASH_DATE = TransformerDate.fit_transform(df.CRASH_DATE)
```

In [10]:

```
ImputerCat_with_Unknown = SimpleImputer(missing_values=np.nan, strategy='constant', fill_value='unknown')
df.loc[:,dfi.NaN_obj_colmns] = ImputerCat_with_Unknown.fit_transform(df.loc[:,dfi.NaN_obj_colmns] )
```

In [11]:

```
ImputerMostFrequent = SimpleImputer(missing_values=np.NaN, strategy='most_frequent')
df.loc[:,dfi.NaN_num_columns] = ImputerMostFrequent.fit_transform(df.loc[:,dfi.NaN_num_columns])
```

FEATURES ENGINEERING

POSTED_SPEED_LIMIT

Acceleration is a major cause of car accidents, which is why we consider it a very important factor for our model.

unfortunately the information provided in the dataset relating to the variable "POSTED_SPEED_LIMIT" is somewhat confused and anomalous for this reason we refer to the official Illinois statute section Vehicle Code no. (625 ILCS 5 /) as a reliable source to sort our data as appropriately as possible.

- 70 mph – Rural interstates and tollways.
- 65 mph – Highways with four lanes.

- 55 mph – Other highways and rural areas.
- 45 mph – Urban freeways and interstates through downtown Chicago.
- 30 mph – Urban Districts – City/Town Areas.
- 20 mph – School Zones (on school days between 7 a.m. and 4 p.m.
- 15 mph – Urban Alleys.

For more info: [Illinois Compiled Statutes](#)

```
In [12]: df.POSTED_SPEED_LIMIT.unique()

Out[12]: array([30, 20, 10, 15, 25, 35, 55, 45,  5, 50,  0, 40, 39, 65,  1, 34,  3,
        12, 60,  9, 26, 33, 49, 70, 11, 36, 32, 24,  2, 14, 31, 29])

In [13]: TransformSpeed = SpeedTransformer()
df['POSTED_SPEED_LIMIT'] = TransformSpeed.fit_transform(df['POSTED_SPEED_LIMIT'])

In [14]: df.POSTED_SPEED_LIMIT.sort_values().unique()

Out[14]: array(['15', '20', '30', '45', '55', '65', '70'], dtype=object)
```

PRIM CONTRIBUTORY CAUSE

```
In [15]: df.PRIM_CONTRIBUTORY_CAUSE.unique()

Out[15]: array(['FAILING TO YIELD RIGHT-OF-WAY', 'UNABLE TO DETERMINE',
        'OPERATING VEHICLE IN ERRATIC, RECKLESS, CARELESS, NEGLIGENT OR AGGRESSIVE MANNER',
        'IMPROPER BACKING', 'IMPROPER OVERTAKING/PASSING',
        'IMPROPER LANE USAGE', 'DISREGARDING TRAFFIC SIGNALS',
        'NOT APPLICABLE', 'IMPROPER TURNING/NO SIGNAL',
        'FOLLOWING TOO CLOSELY', 'DRIVING SKILLS/KNOWLEDGE/EXPERIENCE',
        'ROAD CONSTRUCTION/MAINTENANCE', 'WEATHER',
        'PHYSICAL CONDITION OF DRIVER',
        'FAILING TO REDUCE SPEED TO AVOID CRASH',
        'EVASIVE ACTION DUE TO ANIMAL, OBJECT, NONMOTORIST',
        'HAD BEEN DRINKING (USE WHEN ARREST IS NOT MADE)',
        'DISTRACTION – FROM INSIDE VEHICLE',
        'DRIVING ON WRONG SIDE/WRONG WAY',
        'VISION OBSCURED (SIGNS, TREE LIMBS, BUILDINGS, ETC.)',
        'EQUIPMENT – VEHICLE CONDITION', 'DISREGARDING STOP SIGN',
        'TEXTING', 'DISTRACTION – FROM OUTSIDE VEHICLE',
        'DISREGARDING YIELD SIGN', 'CELL PHONE USE OTHER THAN TEXTING',
        'UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN ARREST IS EFFECTED)',
        'TURNING RIGHT ON RED', 'DISREGARDING ROAD MARKINGS',
        'DISREGARDING OTHER TRAFFIC SIGNS', 'ANIMAL',
        'RELATED TO BUS STOP', 'ROAD ENGINEERING/SURFACE/MARKING DEFECTS',
        'OBSTRUCTED CROSSWALKS',
        'DISTRACTION – OTHER ELECTRONIC DEVICE (NAVIGATION DEVICE, DVD PLAYER, ETC.)',
        'PASSING STOPPED SCHOOL BUS',
        'BICYCLE ADVANCING LEGALLY ON RED LIGHT',
        'MOTORCYCLE ADVANCING LEGALLY ON RED LIGHT'], dtype=object)

In [16]: df.PRIM_CONTRIBUTORY_CAUSE = Target_Grouper().fit_transform(df.PRIM_CONTRIBUTORY_CAUSE)

In [17]: df.PRIM_CONTRIBUTORY_CAUSE.unique()

Out[17]: array(['FAILING TO YIELD RIGHT-OF-WAY', 'UNABLE TO DETERMINE',
        'OPERATING VEHICLE IN ERRATIC, RECKLESS, CARELESS, NEGLIGENT OR AGGRESSIVE MANNER',
        'IMPROPER MANEUVER', 'DISREGARDING', 'NOT APPLICABLE',
        'DRIVING SKILLS/KNOWLEDGE/EXPERIENCE', 'ROAD_CONDITION', 'WEATHER',
        'PHYSICAL CONDITION OF DRIVER',
        'FAILING TO REDUCE SPEED TO AVOID CRASH', 'ANIMAL',
        'ALCOHOL_DRUGS', 'DISTRACTION', 'EQUIPMENT – VEHICLE CONDITION',
        'TURNING RIGHT ON RED', 'RELATED TO BUS STOP',
        'PASSING STOPPED SCHOOL BUS',
        'BICYCLE ADVANCING LEGALLY ON RED LIGHT',
        'MOTORCYCLE ADVANCING LEGALLY ON RED LIGHT'], dtype=object)
```

```
In [18]: df[sorted(set(df.columns)-excluded_cols)].info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 92092 entries, 0 to 92091
Data columns (total 35 columns):
#      Column                                Non-Null Count  Dtype
---  -
0      ALIGNMENT                                92092 non-null  object
1      CRASH_DATE                               92092 non-null  datetime64[ns]
2      CRASH_DAY_OF_WEEK                        92092 non-null  int64
3      CRASH_HOUR                               92092 non-null  int64
4      CRASH_MONTH                              92092 non-null  int64
5      DEVICE_CONDITION                         92092 non-null  object
6      DOORING_I                                92092 non-null  object
7      FIRST_CRASH_TYPE                         92092 non-null  object
8      HIT_AND_RUN_I                            92092 non-null  object
9      INJURIES_FATAL                           92092 non-null  float64
10     INJURIES_INCAPACITATING                  92092 non-null  float64
11     INJURIES_NON_INCAPACITATING              92092 non-null  float64
12     INJURIES_NO_INDICATION                   92092 non-null  float64
13     INJURIES_REPORTED_NOT_EVIDENT            92092 non-null  float64
14     INJURIES_TOTAL                           92092 non-null  float64
15     INJURIES_UNKNOWN                         92092 non-null  float64
16     INTERSECTION_RELATED_I                  92092 non-null  object
```

```
17  LATITUDE                92092 non-null float64
18  LIGHTING_CONDITION      92092 non-null object
19  LONGITUDE               92092 non-null float64
20  MOST_SEVERE_INJURY      92092 non-null object
21  NOT_RIGHT_OF_WAY_I      92092 non-null object
22  NUM_UNITS               92092 non-null int64
23  POSTED_SPEED_LIMIT      92092 non-null object
24  PRIM_CONTRIBUTORY_CAUSE  92092 non-null object
25  ROADWAY_SURFACE_COND    92092 non-null object
26  ROAD_DEFECT             92092 non-null object
27  STREET_DIRECTION        92092 non-null object
28  STREET_NAME             92092 non-null object
29  STREET_NO               92092 non-null int64
30  TRAFFICWAY_TYPE         92092 non-null object
31  TRAFFIC_CONTROL_DEVICE  92092 non-null object
32  WEATHER_CONDITION       92092 non-null object
33  WORKERS_PRESENT_I       92092 non-null object
34  WORK_ZONE_I             92092 non-null object
dtypes: datetime64[ns](1), float64(9), int64(5), object(20)
memory usage: 24.6+ MB
```

```
In [19]: statistic_cols.update(['CRASH_DATE','CRASH_DAY_OF_WEEK','CRASH_HOUR','CRASH_MONTH', 'DOORING_I',
                              'HIT_AND_RUN_I','INJURIES_FATAL','INJURIES_INCAPACITATING',
                              'INJURIES_NON_INCAPACITATING','INJURIES_NO_INDICATION','INJURIES_REPORTED_NOT_EVIDENT',
                              'INJURIES_TOTAL','INJURIES_UNKNOWN','LATITUDE','LONGITUDE','MOST_SEVERE_INJURY',
                              'NUM_UNITS','STREET_DIRECTION','STREET_NAME','STREET_NO'])
```

```
In [20]: features_cols.update(set(df.columns)-excluded_cols.union(statistic_cols))
```

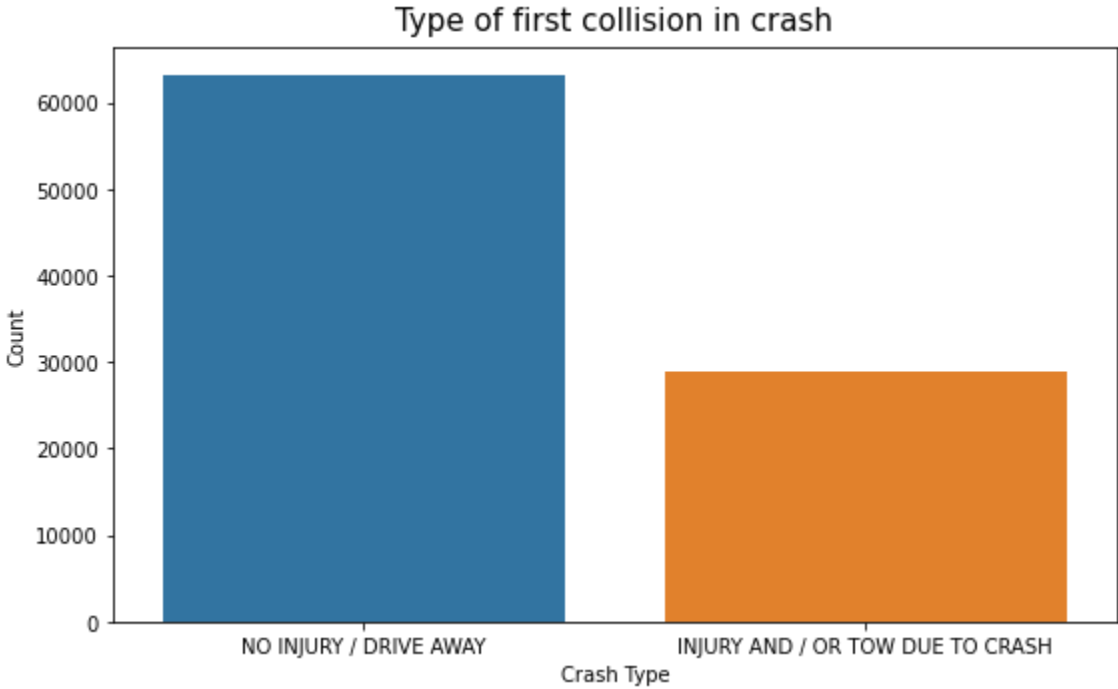
```
In [21]: df[sorted(features_cols)].info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 92092 entries, 0 to 92091
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ALIGNMENT                            92092 non-null  object
1   DEVICE_CONDITION                     92092 non-null  object
2   FIRST_CRASH_TYPE                     92092 non-null  object
3   INTERSECTION_RELATED_I              92092 non-null  object
4   LIGHTING_CONDITION                  92092 non-null  object
5   NOT_RIGHT_OF_WAY_I                  92092 non-null  object
6   POSTED_SPEED_LIMIT                  92092 non-null  object
7   PRIM_CONTRIBUTORY_CAUSE             92092 non-null  object
8   ROADWAY_SURFACE_COND                92092 non-null  object
9   ROAD_DEFECT                         92092 non-null  object
10  TRAFFICWAY_TYPE                     92092 non-null  object
11  TRAFFIC_CONTROL_DEVICE               92092 non-null  object
12  WEATHER_CONDITION                   92092 non-null  object
13  WORKERS_PRESENT_I                   92092 non-null  object
14  WORK_ZONE_I                         92092 non-null  object
dtypes: object(15)
memory usage: 10.5+ MB
```

EXPLORATORY DATA ANALYSIS

CRASH TYPE

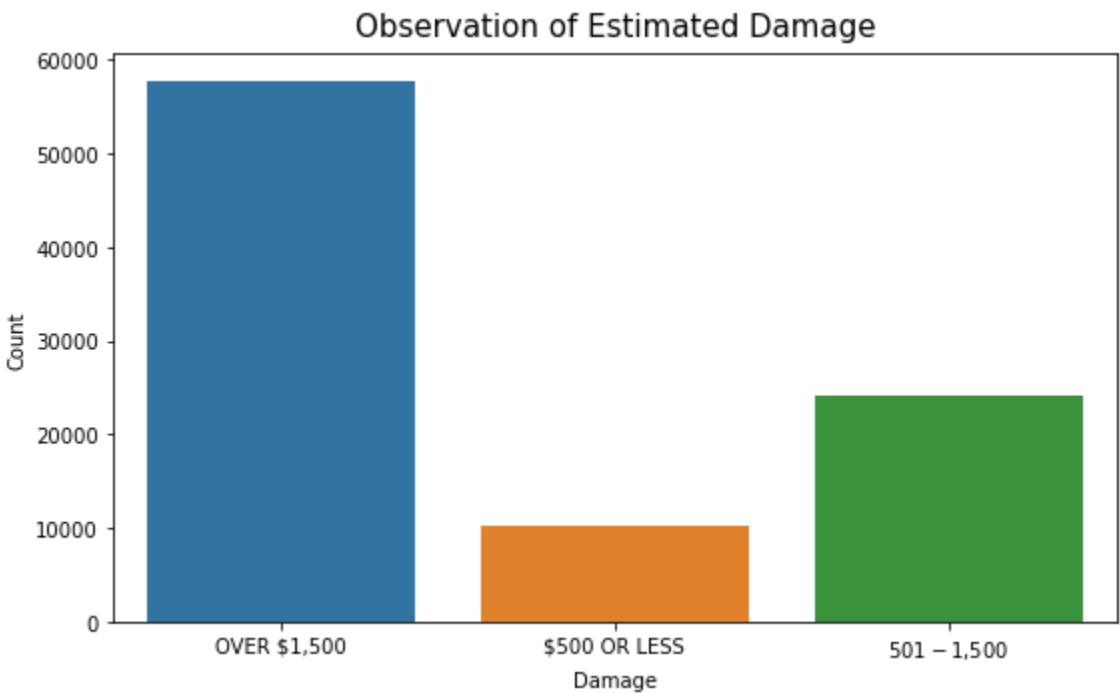
```
In [19]: fig, ax = plt.subplots(figsize=(8,5))
sns_crach_type = sns.countplot(x=df.CRASH_TYPE, ax=ax)
ax.set_title("Type of first collision in crash", fontsize=15, ha='center', va='bottom')
ax.set_xlabel('Crash Type')
ax.set_ylabel('Count')
plt.savefig('./img/crach_type.jpg',format="jpg")
plt.show()
```



ESTIMATED DAMAGE

```
In [20]: fig, ax = plt.subplots(figsize=(8,5))
sns_crach_type = sns.countplot(x=df.DAMAGE, ax=ax)
ax.set_title("Observation of Estimated Damage", fontsize=15, ha='center', va='bottom')
```

```
ax.set_xlabel('Damage')
ax.set_ylabel('Count')
plt.savefig('./img/damage.jpg',format="jpg")
plt.show()
```

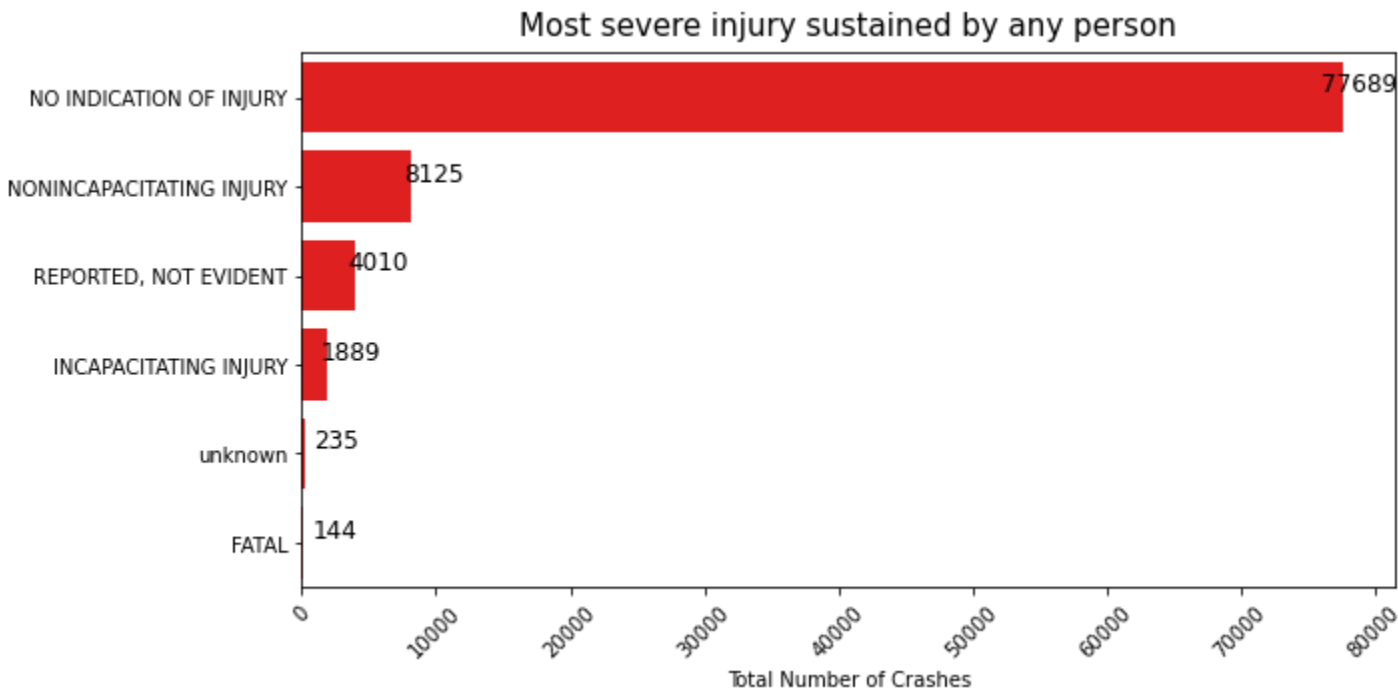


MOST SEVERE INJURY

```
In [24]: fig, ax = plt.subplots(figsize=(10,5))
msi = sns.barplot(x=df.MOST_SEVERE_INJURY.value_counts().values,
                  y=df.MOST_SEVERE_INJURY.value_counts().index, color='red', ax=ax)

for index, row in df.MOST_SEVERE_INJURY.value_counts().iteritems():
    msi.text(x=4000+row, y=df.MOST_SEVERE_INJURY.value_counts().index.get_loc(index),
             s=round(row), color='black', va='bottom', ha="right", fontsize=12)

ax.set_title("Most severe injury sustained by any person", fontsize=15, ha='center', va='bottom')
ax.set_xlabel('Total Number of Crashes')
# ax.set_ylabel('Count')
plt.xticks(rotation=45)
fig.tight_layout()
plt.savefig('./img/Most_Severe_Injury.jpg', dpi=150, format="jpg")
plt.show()
```




WEATHER CONDITION

```
In [28]: fig, ax = plt.subplots(figsize=(10,5))
wc = sns.barplot(x=df.WEATHER_CONDITION.value_counts().values,
                 y=df.WEATHER_CONDITION.value_counts().index, color='c', ax=ax)


for index, row in df.WEATHER_CONDITION.value_counts().iteritems():
    wc.text(x=row,
            y=df.WEATHER_CONDITION.value_counts().index.get_loc(index),
            s='{0:.2f}%'.format(round((row/df.shape[0]*100),2)),
            color='black', fontsize=10, rotation_mode='anchor',
            verticalalignment='center_baseline', horizontalalignment='left')

ax.set_title("Weather condition at time of crash", fontsize=15, ha='center', va='bottom')
ax.set_xlabel('Total Number of Crashes')
plt.savefig('./img/Weathercondition.jpg',format="jpg")
plt.show()
```


ROAD SURFACE CONDITION

```
In [93]: fig, ax = plt.subplots(figsize=(8,5))
sns_crach_type = sns.countplot(x=df.ROADWAY_SURFACE_COND, ax=ax)
ax.set_title("Road surface condition", fontsize=15, ha='center', va='bottom')
ax.set_xlabel('Road surface condition')
ax.set_ylabel('Count')
plt.xticks(rotation=45)
plt.savefig('./img/road_surface_condition.jpg', dpi=150, format="jpg")
plt.show()
```



Posted Speed Limit Analysis


```
In [86]: fig, ax = plt.subplots(figsize=(8,5))
sns.barplot(x=df.POSTED_SPEED_LIMIT.value_counts().index, y=df.POSTED_SPEED_LIMIT.value_counts(), ax=ax)
ax.set_title("Car Crashes Vs. Posted Speed Limit", fontsize=15, ha='center', va='bottom')
ax.set_xlabel('Posted Speed Limit')
ax.set_ylabel('Number of crashes')
plt.savefig('./img/speed_number_of_crashes.jpg', format="jpg")
plt.show()
```



Speed Limit by Number of units-related

```
In [87]: speed_unit = df.groupby('POSTED_SPEED_LIMIT', as_index=False).agg({'NUM_UNITS': 'sum'})

fig, ax = plt.subplots(figsize=(8,5))
sns_speed_unit=sns.barplot(x='POSTED_SPEED_LIMIT', y='NUM_UNITS', data=speed_unit, ax=ax)
for index, row in speed_unit.iterrows():
    sns_speed_unit.text(x=row.name,
                        y=row.NUM_UNITS,
                        s=round(row.NUM_UNITS),
                        color='black', va='bottom', ha="center", fontsize=10)
ax.set_title("Number of units involved in the crash", fontsize=15, ha='center', va='bottom')
ax.set_xlabel('Posted Speed Limit')
ax.set_ylabel('Number of Units')
plt.savefig('./img/speed_unit.jpg',format="jpg")
plt.show()
```




INJURIES ANALYSIS

Number of injuries by Speed Limit-related

```
In [88]: speed_injuries = df.groupby('POSTED_SPEED_LIMIT', as_index=False).agg({'INJURIES_TOTAL': 'sum'})

fig, ax = plt.subplots(figsize=(8,5))
sns_speed_injured=sns.barplot(x='POSTED_SPEED_LIMIT', y='INJURIES_TOTAL', data=speed_injuries, ax=ax)
for index, row in speed_injuries.iterrows():
    sns_speed_injured.text(x=row.name,
                           y=row.INJURIES_TOTAL,
                           s=round(row.INJURIES_TOTAL),
                           color='black', va='bottom', ha="center", fontsize=10)
ax.set_title("Number of injuries related Speed Limit", fontsize=15, ha='center', va='bottom')
ax.set_xlabel('Posted Speed Limit')
ax.set_ylabel('Total Injuries')
plt.savefig('./img/speed_injures.jpg',format="jpg")
plt.show()
```




Number of injuries by Month-related

```
In [89]: injuries_month = df.groupby('CRASH_MONTH', as_index=False).agg({'INJURIES_TOTAL': 'sum'})


fig, ax = plt.subplots(figsize=(8,5))
sns_injuries_month=sns.barplot(x='CRASH_MONTH', y='INJURIES_TOTAL', data=injuries_month, ax=ax)
for index, row in injuries_month.iterrows():
    sns_injuries_month.text(x=row.name,
                             y=row.INJURIES_TOTAL,
                             s=round(row.INJURIES_TOTAL),
                             color='black', va='bottom', ha="center", fontsize=10)
ax.set_title("The month component of CRASH_DATE", fontsize=15, ha='center', va='bottom')
ax.set_xlabel('Month')
```

```
ax.set_ylabel('Total Injuries')
plt.savefig('./img/injuries_month.jpg',format="jpg")
plt.show()
```



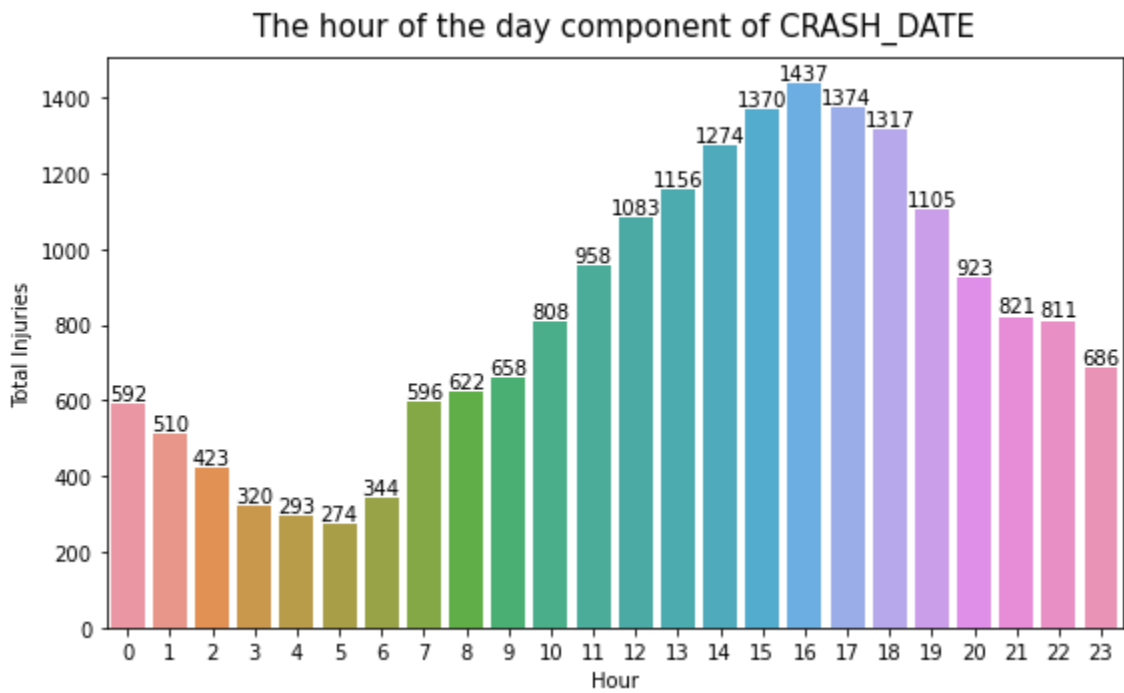
Number of injuries by Day of Week-related

```
In [90]: injuries_week = df.groupby('CRASH_DAY_OF_WEEK', as_index=False).agg({'INJURIES_TOTAL': 'sum'})
fig, ax = plt.subplots(figsize=(8,5))
sns_injuries_week=sns.barplot(x='CRASH_DAY_OF_WEEK', y='INJURIES_TOTAL', data=injuries_week, ax=ax)
for index, row in injuries_week.iterrows():
    sns_injuries_week.text(x=row.name,
                           y=row.INJURIES_TOTAL,
                           s=round(row.INJURIES_TOTAL),
                           color='black', va='bottom', ha="center", fontsize=10)
ax.set_title("The day of the week component of CRASH_DATE", fontsize=15, ha='center', va='bottom')
ax.set_xlabel('Week')
ax.set_ylabel('Total Injuries')
plt.savefig('./img/injuries_week.jpg',format="jpg")
plt.show()
```



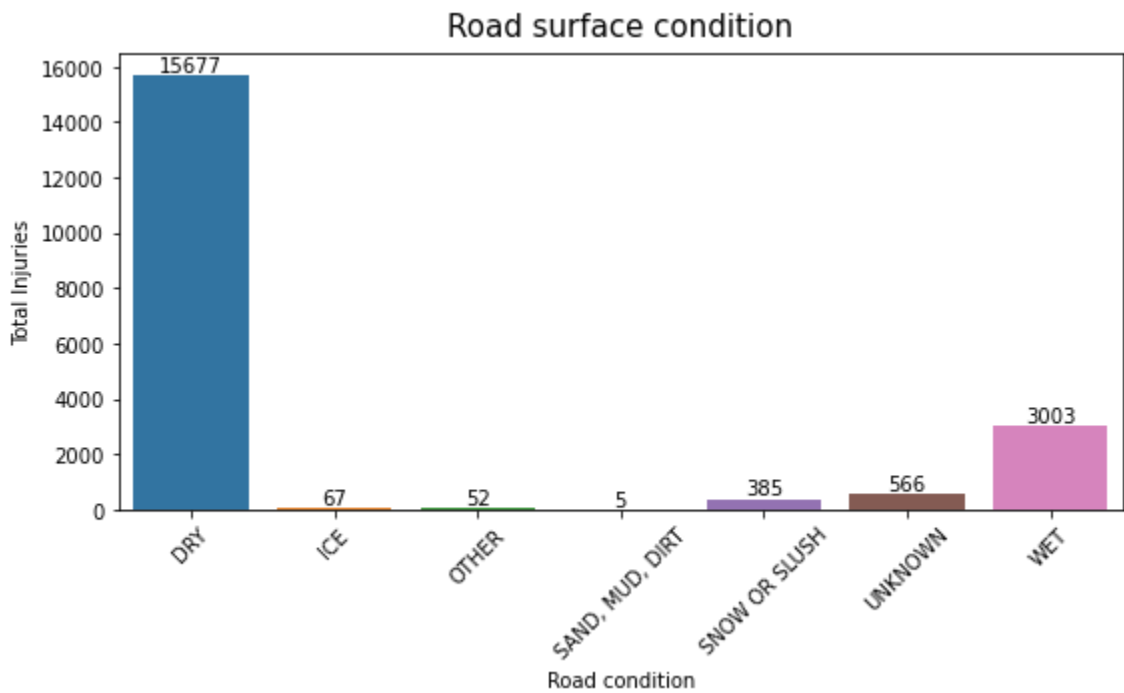
Number of injuries by Hour-related

```
In [91]: injuries_hour = df.groupby('CRASH_HOUR', as_index=False).agg({'INJURIES_TOTAL': 'sum'})
fig, ax = plt.subplots(figsize=(8,5))
sns_injuries_hour=sns.barplot(x='CRASH_HOUR', y='INJURIES_TOTAL', data=injuries_hour, ax=ax)
for index, row in injuries_hour.iterrows():
    sns_injuries_hour.text(x=row.name,
                           y=row.INJURIES_TOTAL,
                           s=round(row.INJURIES_TOTAL),
                           color='black', va='bottom', ha="center", fontsize=10)
ax.set_title("The hour of the day component of CRASH_DATE", fontsize=15, ha='center', va='bottom')
ax.set_xlabel('Hour')
ax.set_ylabel('Total Injuries')
plt.savefig('./img/injuries_hour.jpg',format="jpg")
plt.show()
```

Number of injuries by Roadway Condition-related

```
In [92]: injuries_road = df.groupby('ROADWAY_SURFACE_COND', as_index=False).agg({'INJURIES_TOTAL': 'sum'})
fig, ax = plt.subplots(figsize=(8,5))
sns_injuries_road=sns.barplot(x='ROADWAY_SURFACE_COND', y='INJURIES_TOTAL', data=injuries_road, ax=ax)
for index, row in injuries_road.iterrows():
    sns_injuries_road.text(x=row.name,
                           y=row.INJURIES_TOTAL,
                           s=round(row.INJURIES_TOTAL),
                           color='black', va='bottom', ha="center", fontsize=10)
sns.barplot(x='ROADWAY_SURFACE_COND', y='INJURIES_TOTAL', data=injuries_road, ax=ax)
plt.xticks(rotation=45)
ax.set_title("Road surface condition", fontsize=15, ha='center', va='bottom')
ax.set_xlabel('Road condition')
ax.set_ylabel('Total Injuries')
plt.savefig('./img/Road_surface_condition.jpg', format="jpg")
plt.show()
```



PREPROCESSING

```
In [35]: df_features = df[sorted(features_cols)].copy()
```

```
In [36]: df_features_info = Dframe(df_features)
```

```
In [37]: df_features.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 92092 entries, 0 to 92091
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ALIGNMENT                            92092 non-null  object
1   DEVICE_CONDITION                     92092 non-null  object
2   FIRST_CRASH_TYPE                    92092 non-null  object
3   INTERSECTION_RELATED_I              92092 non-null  object
4   LIGHTING_CONDITION                  92092 non-null  object
5   NOT_RIGHT_OF_WAY_I                  92092 non-null  object
6   POSTED_SPEED_LIMIT                  92092 non-null  object
7   PRIM_CONTRIBUTORY_CAUSE              92092 non-null  object
8   ROADWAY_SURFACE_COND                 92092 non-null  object
9   ROAD_DEFECT                          92092 non-null  object
10  TRAFFICWAY_TYPE                      92092 non-null  object
11  TRAFFIC_CONTROL_DEVICE               92092 non-null  object
12  WEATHER_CONDITION                    92092 non-null  object
13  WORKERS_PRESENT_I                    92092 non-null  object
14  WORK_ZONE_I                          92092 non-null  object
dtypes: object(15)
memory usage: 10.5+ MB
```

```
In [38]: category_cols.update(('TRAFFIC_CONTROL_DEVICE', 'DEVICE_CONDITION', 'WEATHER_CONDITION', 'LIGHTING_CONDITION',
                             'FIRST_CRASH_TYPE', 'TRAFFICWAY_TYPE', 'ALIGNMENT', 'ROADWAY_SURFACE_COND', 'ROAD_DEFECT',
                             'INTERSECTION_RELATED_I', 'NOT_RIGHT_OF_WAY_I', 'WORK_ZONE_I', 'WORKERS_PRESENT_I', 'POSTED_SPEED_LIM
```

Features Transformation using OneHotEncoder

```
In [39]: ohe_features = OneHotEncoder(sparse=False, drop='first')
ohe_features.fit(df_features[category_cols])

df_features_ohe = pd.concat([df_features.drop(labels=category_cols, axis=1),
                             pd.DataFrame(ohe_features.transform(df_features[category_cols]),
                                           columns=ohe_features.get_feature_names(list(category_cols))),
                             axis=1)
df_features_ohe.drop(labels=['PRIM_CONTRIBUTORY_CAUSE'], axis=1, inplace=True)
```

Target Transformation using OneHotEncoder

```
In [40]: ohe_target = OneHotEncoder(handle_unknown='ignore', sparse=False)
df_target_ohe = pd.DataFrame(ohe_target.fit_transform(df_features.PRIM_CONTRIBUTORY_CAUSE.values.reshape(-1,1)),
                             columns=ohe_target.categories_)
df_target_ohe.columns = df_target_ohe.columns.get_level_values(0).str.replace(' ', '_')

In [41]: lc_prim_contributory_cause = LabelEncoder()
df_target_encoded = pd.DataFrame(lc_prim_contributory_cause.fit_transform(df_features['PRIM_CONTRIBUTORY_CAUSE']),
                                 columns=['PRIM_CONTRIBUTORY_CAUSE'])

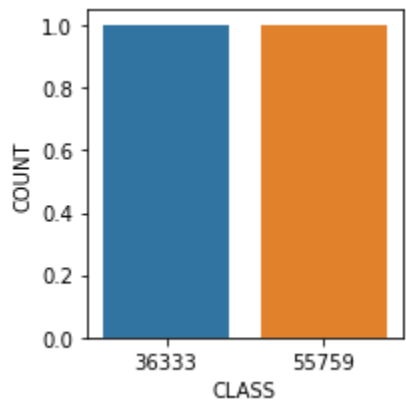
y_encoded_dict = dict(zip(lc_prim_contributory_cause.classes_,
                          lc_prim_contributory_cause.transform(lc_prim_contributory_cause.classes_)))
```

MODELING

it's broad clear we're dealing with a multi-class & multi-label imbalanced dataset with 99% of categorical data except for POSTED_SPEED_LIMIT. we end up with 102 features our goal is to validate one of the main features of car crashes which is "UNABLE_TO_DETERMINE", that rappresents about 40% of "PRIM_CONTRIBUTORY_CAUSE". the dataset.

```
In [42]: fig, ax = plt.subplots(figsize=(3,3))
sns.countplot(x=df_target_ohe.UNABLE_TO_DETERMINE.value_counts())
ax.set_xlabel('CLASS')
ax.set_ylabel('COUNT')
```

Out[42]: Text(0, 0.5, 'COUNT')



FEATURE SELECTION

Since we're dealing with a remarkably large dataset with 102 features, we're going to process our features by removing irrelevant and redundant data in order to reduce computation time and improve learning accuracy.

```
In [432... alphas = np.logspace(-3, -1, 100)
lasso = LassoCV(alphas=alphas, max_iter=100_000, n_jobs=6, random_state=264, cv=5)
sfm = SelectFromModel(lasso, threshold=0.0001)
sfm.fit_transform(df_features_ohe, np.ravel(np.array(df_target_ohe['UNABLE_TO_DETERMINE'])))
feature_selection_col = df_features_ohe.columns[sfm.get_support()]
```

```
In [43]: X = df_features_ohe
y = np.ravel(np.array(df_target_ohe['UNABLE_TO_DETERMINE']))
print(f"Type of Target : {type_of_target(y)}")

Type of Target : binary
```

```
In [44]: x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=.25, random_state=264, stratify=y)
x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

Out[44]: ((69069, 107), (23023, 107), (69069,), (23023,))

```
In [112... x_train, y_train = check_X_y(x_train, y_train)
```

Decision Tree

```
In [272... class_weight = get_class_weight(x_train, y_train)

In [273... dc_param={
    'criterion': 'entropy',
    'class_weight': class_weight,
    'random_state': 264 }

In [199... joblib.dump(dc_param, './pkl/DecisionTree_params.pkl')

Out[199... ['./pkl/DecisionTree_params.pkl']]

In [274... cls_tree = DecisionTreeClassifier(**dc_param)

In [278... cls_tree.fit(x_train, y_train)

Out[278... DecisionTreeClassifier(class_weight={0.0: 0.8258088428704656,
                                     1.0: 1.2673211009174312},
                        criterion='entropy', random_state=264)

In [297... model1_y_train_pred = model_metric(cls_tree, x_train, y_train)

ROC_AUC SCORE    0.7376198010166575
ACCURACY SCORE   0.7403610881871752
```

Random Forest

```
In [298... cls_rf = RandomForestClassifier(class_weight=class_weight, n_jobs=6, random_state=264)

In [299... cls_rf.fit(x_train, np.ravel(np.array(y_train)))

Out[299... RandomForestClassifier(class_weight={0.0: 0.8258088428704656,
                                     1.0: 1.2673211009174312},
                        n_jobs=6, random_state=264)

In [300... model2_y_train_pred = model_metric(cls_rf, x_train, y_train)

ROC_AUC SCORE    0.7360246988386605
ACCURACY SCORE   0.7416496546931329
```

RandomForestClassifier: Hyperparameters with GridSearchCV

```
In [146... rfc_params = {
    'n_estimators': [2000],
    'criterion': ['gini', 'entropy'],
    'max_depth': [150],
    'max_features': [30],
    'bootstrap': [True],
    'oob_score': [True],
    'n_jobs': [6],
    'class_weight': [class_weight],
    'random_state': [264]}

In [588... cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=1, random_state=264)
grid_search = GridSearchCV(cls_rf, param_grid=rfc_params, scoring='roc_auc', n_jobs=6, cv=cv)

In [589... grid_search.fit(x_train, np.ravel(np.array(y_train)))

Out[589... GridSearchCV(cv=RepeatedStratifiedKFold(n_repeats=1, n_splits=10, random_state=264),
            estimator=RandomForestClassifier(class_weight={0.0: 0.8258088428704656,
                                                         1.0: 1.2673211009174312},
                                             n_jobs=6, random_state=264),
            n_jobs=6,
            param_grid={'bootstrap': [True],
                        'class_weight': [{0.0: 0.8258088428704656,
                                         1.0: 1.2673211009174312}],
                        'criterion': ['gini', 'entropy'], 'max_depth': [150],
                        'max_features': [30], 'n_estimators': [2000],
                        'n_jobs': [6], 'oob_score': [True],
                        'random_state': [264]}},
            scoring='roc_auc')

In [698... rf_best_params=grid_search.best_params_
joblib.dump(rf_best_params, './pkl/RF_Grid_Search_best_params.pkl', compress=9)

In [148... cls_rf = RandomForestClassifier(**rf_best_params)
model3_train = cls_rf.fit(x_train, y_train)

In [150... model3_y_train_pred = model_metric(cls_rf, x_train, y_train)

ROC_AUC SCORE    0.7365930367018547
ACCURACY SCORE   0.7423301336344814
```

In [309... joblib.dump(model3_y_train_pred, './pk1/model3_y_train_predicted.pkl')

Out[309... ['./pk1/model3_y_train_predicted.pkl']

TARGET PREDICTION

In [307... model3_test = cls_rf.fit(x_test, y_test)
model3_y_test_pred = model_metric(cls_rf, x_test, y_test)

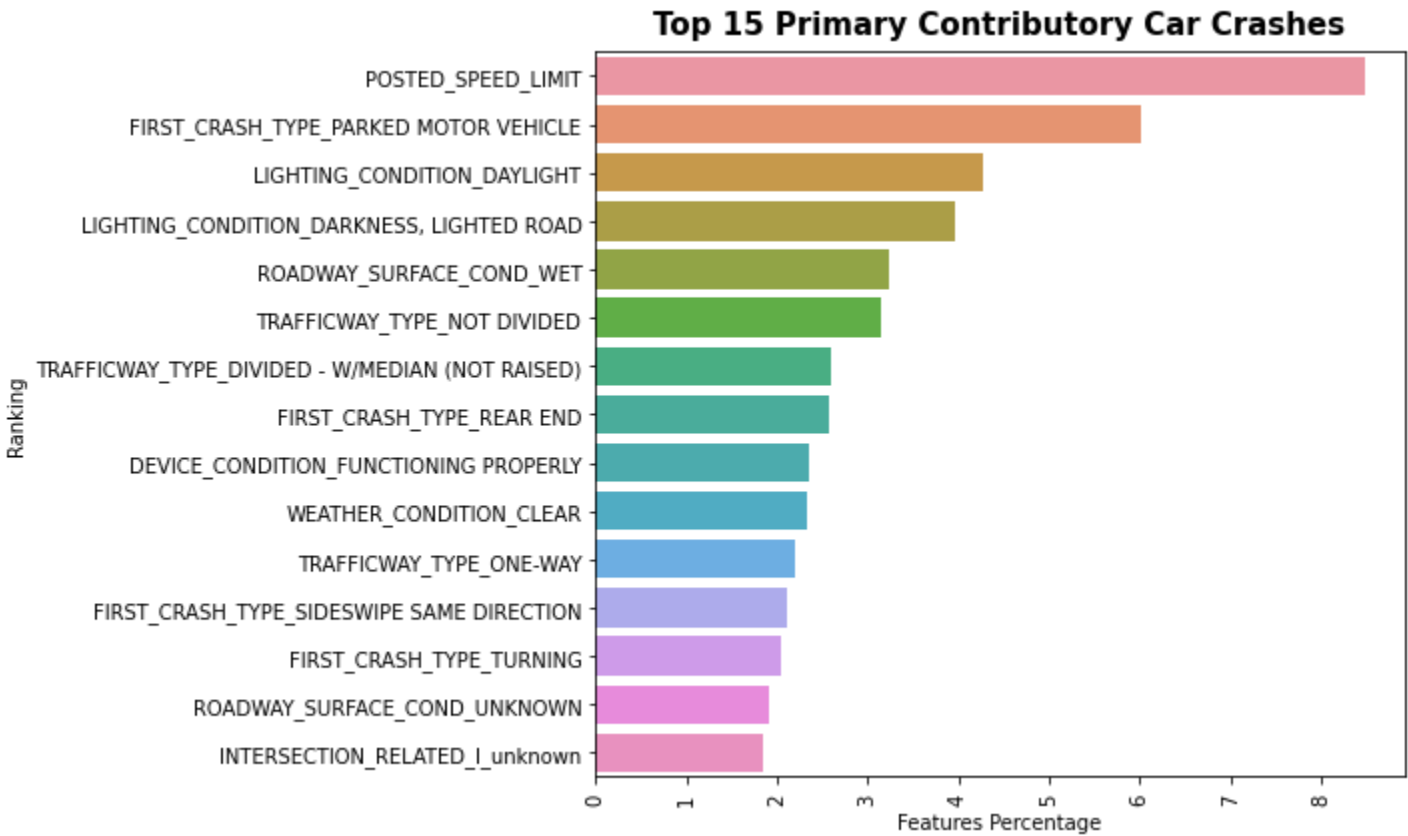
ROC_AUC SCORE 0.7728139273851178
ACCURACY SCORE 0.7803500846979108

In [308... joblib.dump(model3_y_test_pred, './pk1/model3_y_test_predicted.pkl')

Out[308... ['./pk1/model3_y_test_predicted.pkl']

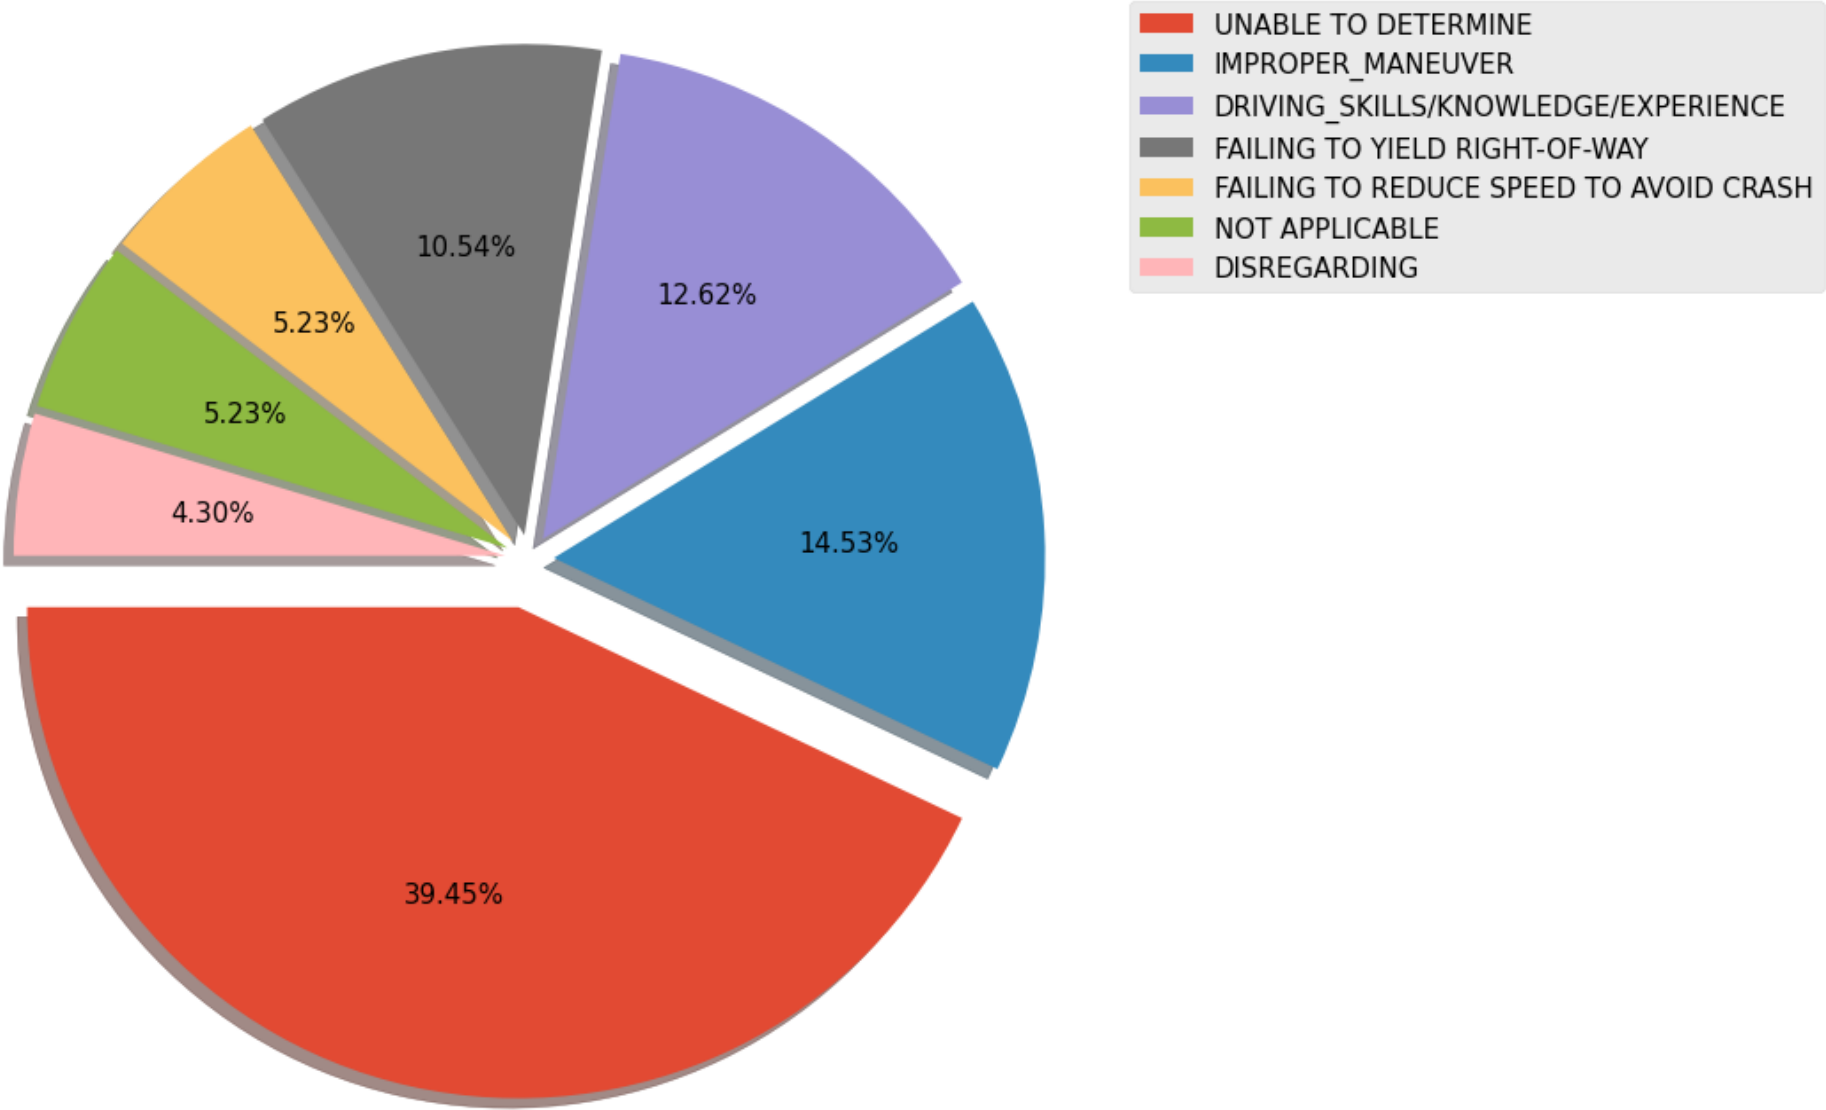
In [152... Features_importance_df = pd.DataFrame({'Features': x_train.columns,
'Ranking':cls_rf.feature_importances_}).sort_values(by='Ranking',ascending=False)
Features_importance_df['Ranking'] = Features_importance_df['Ranking'].apply(lambda x: x*100)

In [154... fig, ax = plt.subplots(figsize=(10,6))
g = sns.barplot(x='Ranking', y='Features', data=Features_importance_df.head(15))
plt.xticks(rotation=90)
ax.set_title("Top 15 Primary Contributory Car Crashes", fontsize=15, weight='bold', ha='center', va='bottom')
ax.set_xlabel('Features Percentage')
ax.set_ylabel('Ranking')
plt.savefig('./img/Y_primary_contributory_crashes.jpg',format="jpg")
plt.show()



In [101... pcc = df_features.PRIM_CONTRIBUTORY_CAUSE.value_counts(normalize=True).apply(lambda x:x*100).sort_values(ascending=False)
pcc=pcc[pcc>=2]
labels = pcc.index
sizes = pcc.values
explode = pcc.apply(lambda x:0.05)
explode[0]=0.1
explode = pcc.apply(lambda x:0.2 if x > 20 else 0)
plt.style.use('ggplot')
fig1, ax1 = plt.subplots(figsize=(15,10))
ax1.pie(x=sizes, explode=explode,
autopct=lambda v: '{p:.2f}%'.format(p=round(v/100*sizes.sum(),2)),
shadow=True,
startangle=180,
textprops={'fontsize': 15})
ax1.axis('equal')
plt.legend(labels, loc='upper left', bbox_to_anchor=(1, 1), fontsize=15)
ax1.set_title("PRIMARY CONTRIBUTORY CAR CRASHES", fontsize=30, ha='center', va='bottom')
plt.savefig('./img/top_15_contributory.jpg', dpi=200, format="jpg")
plt.show()

PRIMARY CONTRIBUTORY CAR CRASHES



Interpreting Data

Based on the data, 40% of accidents in the year 2020 are "UN ABLE TO DETERMINE", followed by "FAILING TO YIELD THE RIGHT-OF-WAY" and "FOLLOWING TO CLOSELY" which represent 10% and 9% respectively.

According to our classification model, the road accidents that seem to occur the most are caused by the speeding where the posted speed limit is 30mph.

Random Forest Model predicts car crashes with 78% accuracy even dealing with a multiclass imbalanced dataset. the accuracy score between the 2 models (training model 74.23% - test model 78.03%) are too close which indicates that we avoided over-fitting